

Assignment 3: Pipeline Automation & Data Validation

Onboarding Flow Optimization A/B Test

Aishwarya Mishra | USN 2648610 | MSc Computational Statistics & Applied AI, Christ University

Table of Contents

Abstract

1. What “Automate the Pipeline” Actually Requires

2. Error Taxonomy: Retry vs. Alert

2.1 Retry Mechanics

3. Logging

4. Validation

5. Monitoring

6. Alerting

7. Scheduling

8. Reliability: Failure Scenarios Considered

9. Assumptions and Limitations

10. Observed Behavior

11. Execution Evidence

11.1 Normal Run

11.2 Transient Failure That Recovers

11.3 Retries Exhausted → Alert

11.4 Fatal Error — No Retry

11.5 Validation Failure — No Retry

11.6 Monitoring: Health Summary Across Runs


11.7 Scheduler Firing Live

12. How to Run

Abstract

Assignment 2 produced a working, correct pipeline that a human runs by typing `python build_db.py`. This assignment removes the human from that loop: the pipeline runs on a schedule, classifies its own failures, retries the ones worth retrying, alerts an owner on the ones that aren't, validates its own output against a policy before calling itself successful, and leaves behind a queryable history of every run so a person (or another automated check) can tell how the pipeline has been doing without reading raw logs. None of Assignment 2's SQL changes; this assignment wraps it.

1. What “Automate the Pipeline” Actually Requires



Pipeline automation and reliability flow

Figure 1. Control flow through `pipeline_runner.py`, from the scheduler trigger to the final alert or monitoring record. The left branch is the retry loop (transient errors); the right branch is the immediate-alert path (fatal errors); the center branch is the success/validation path.

Turning a manually-run script into an unattended pipeline is not just “put it in cron.” A script that only ever ran once, by a person watching the terminal, gets four things for free that an unattended job does not: a human notices if it didn’t finish, a human notices if the output looks wrong, a human can just re-run it if something transient went wrong, and a human remembers roughly how it’s been going lately. Automation has to replace all four of those, explicitly:

What a human provided	What replaces it here
Notices the job didn't finish	Structured logging (Section 3) + a rotating log file
Notices the output looks wrong	An explicit, executable validation policy (Section 4) run after every build
Re-runs it if something transient failed	Automatic retry with backoff, but <i>only</i> for errors classified as worth retrying (Section 2)
Remembers how it's been going	A monitoring history file and a summary tool (Section 5)
Escalates if it's actually broken	A pluggable Notifier invoked on any terminal failure (Section 6)

2. Error Taxonomy: Retry vs. Alert

The assignment brief's own example — “if there is a connection error or server error, retry after 5 minutes... or notify the owner” — already implies a classification the naive version of automation skips: **not every failure should be retried**. Retrying a failure caused by bad input data just re-produces the same bad output five times before giving up, burning the retry budget and delaying the alert for no benefit. This pipeline draws the line explicitly with two exception types (`pipeline_runner.py`):

	<code>TransientError</code>	<code>FatalError</code>
Meaning	The same operation might succeed with no other change	The operation will fail identically until something is fixed
Examples in this pipeline	SQLite <code>OperationalError</code> (lock contention, momentary I/O issue)	Missing source CSV; malformed <code>stepmap.json</code> ; a policed data-quality check failing
Handled by	<code>run_with_retries()</code> — sleep and retry, up to <code>--max-retries</code>	Propagates immediately to the alerting path, no retry attempted
Why	The pipeline's full-rebuild design (Assignment 2, Section 4) makes a clean retry safe and cheap	Retrying costs time and retry budget while producing an identical failure and an identical, uninformative alert

An exception that doesn't match either known type is treated as **fatal by default**, not silently retried — an unclassified error might indicate a bug in the pipeline itself, and retrying blindly could mask that rather than surface it. The taxonomy is meant to grow: `pipeline_runner.py` 's `except Exception` fallback logs and alerts with an explicit note that the error type should be classified, rather than pretending to have handled it correctly.

2.1 Retry Mechanics

`run_with_retries()` implements the “retry after 5 minutes... limit of 5 retries” policy from the brief directly: `--max-retries` (default 5) and `--retry-delay-seconds` (default 300 = 5 minutes) are both configurable, and a `--backoff` multiplier (default 1.0, fixed-interval) can be set above 1.0 for exponential backoff (5 min, 10 min, 20 min, ...) if a transient condition is more likely to need longer to clear than a fixed

interval assumes. Each retry attempt re-runs the *entire* build from scratch, not just the failed step — safe and correct specifically because Assignment 2’s ingestion is idempotent (drop-and-rebuild), so a partial prior attempt leaves no state a retry needs to clean up first.

3. Logging

`pipeline_runner.py` configures Python’s standard `logging` module once per run, with two handlers sharing one format (`timestamp | level | run=<run_id> | message`):

- A `RotatingFileHandler` (`logs/pipeline.log`, 5 MB per file, 5 files retained) — durable, doesn’t grow unbounded, survives past the process exiting.
- A console `StreamHandler` — visible immediately in an interactive run or in cron’s own captured output.

Every log line carries a `run_id` (a timestamp plus a short random suffix), so a specific run’s complete history can be isolated with `grep run=<run_id> logs/pipeline.log` even though all runs share the same rotating file — chosen over one-file-per-run specifically so log volume stays bounded regardless of how many times the job has fired.

4. Validation

Reuses Assignment 2’s own data-quality checks (`schema/05_data_quality_checks.sql`) rather than inventing a second, possibly-inconsistent set — but adds something Assignment 2 did not have: an explicit, executable **pass/fail policy** (`VALIDATION_POLICY` in `pipeline_runner.py`) that decides which check results are acceptable. Assignment 2 printed twelve numbers for a human to read; Assignment 3 decides, in code, whether those numbers mean the run succeeded.

Eleven of the twelve checks are policed (must equal an expected value, or satisfy a bound); `staging_rejects` is deliberately **not** policed, because a non-zero count there is expected and desired (Assignment 2, Section 5.1 — it’s how the validation logic’s own effectiveness gets verified). Encoding that distinction as a real policy, not a comment, means a genuine regression — say, a future schema change that silently lets `orphan_fact_users` go non-zero — fails the pipeline run loudly instead of printing an ignored number.

A validation failure raises `ValidationFailure` (a subtype of `FatalError`, Section 2): the same input data would fail the same way on a retry, so it goes straight to alerting rather than into the retry loop.

5. Monitoring

Every run — success or failure, however it failed — appends one JSON record to `logs/run_history.jsonl`: `run_id`, start/end timestamps, duration, `status`, `retries_used`, the validation summary (on success), and the error (on failure). This is deliberately separate from the logging file in Section 3: logging answers “what happened inside this one run, step by step” (for debugging one incident); monitoring answers “how has the pipeline been doing across many runs” (for a health check that doesn’t want to parse prose log lines).

`monitor.py` reads that history and reports: a table of recent runs, the success rate over a configurable window, and average duration of successful runs. Its `--check` flag exits non-zero if the *most recent* run failed, or if the success rate over the last five-plus runs drops below 70% even when the latest run happened to pass — catching a pipeline that’s flapping (intermittently failing) and not just one that’s currently down. That exit code

is designed to be chained: a second cron entry (Section 7) runs `monitor.py --check` independently of `pipeline_runner.py`'s own alerting, so a bug in the runner's alerting path doesn't silently remove the *only* signal that something is wrong.

6. Alerting

`notifier.py` defines a `Notifier` interface with one method, `notify(subject, message, severity)`. The default, `LogNotifier`, writes the alert through the same structured logger as everything else (at `CRITICAL`) and to stdout — meaning it works with zero external configuration, which matters for a coursework deliverable graded by someone who won't have SMTP credentials or a Slack workspace to hand. `EmailNotifier` and `SlackWebhookNotifier` implement the same interface against real channels, shown to demonstrate the extension point rather than wired to live credentials; swapping one in for production is a one-line change (`notifier = EmailNotifier(...)` in `main()`), not a redesign. `MultiNotifier` fans one alert out to several channels at once, for a real deployment that wants both a log entry and a page.

`pipeline_runner.py` calls the notifier exactly once per failed run, with a message specific to *why* it failed (which check violated policy; which file was missing; how many retries were exhausted) — an alert that just says “pipeline failed” sends the on-call person straight to the logs anyway, so the alert itself carries the diagnosis where it's known at alert time.

7. Scheduling

Two scheduling mechanisms are provided, aimed at two different needs:

- `scheduler/crontab.txt` — the real answer for a production or graded-as-deployed system: a standard cron entry running `pipeline_runner.py` daily, plus a second, offset entry running `monitor.py --check` independently (Section 5). Documented with the caveat that actually matters in practice: cron's environment is minimal (no `PATH`, no shell profile), so every path in the crontab entry is absolute — a script that works fine when run by hand and then mysteriously “does nothing” under cron is almost always this.
- `scheduler/scheduler_demo.py` — a pure-Python scheduler (the `schedule` library) for environments without OS cron access, or for demonstrating the schedule actually firing without waiting until 2 AM: `--demo-interval-seconds 15` runs the job every fifteen seconds instead of daily, so the retry/validation/alerting behavior in Sections 2–6 can be watched live in one terminal.

8. Reliability: Failure Scenarios Considered

The assignment brief's own example (connection/server error → retry, then alert after N failures) is one instance of a broader question: *what can actually go wrong, and for each, is retrying the right response?* This table is the “think more cases like this” the brief asks for.

Scenario	Classification	Why
Database locked / momentary I/O error	Transient — retry	SQLite lock contention typically clears within seconds; the full-rebuild design makes a retry safe

Scenario	Classification	Why
Simulated “connection/server error” (the brief’s own example)	Transient — retry, then alert after <code>--max-retries</code>	Directly implements the brief’s example via <code>--simulate always-transient</code>
Source CSV missing	Fatal — alert immediately, no retry	Will not resolve itself; retrying just delays discovering the upstream job didn’t produce its output
<code>stepmap.json</code> missing or malformed	Fatal — alert immediately	A config/code error, not an environmental blip
Data-quality policy violated (e.g. <code>mart_kpi_group_count != 3</code>)	Fatal (<code>ValidationFailure</code>) — alert immediately, no retry	Re-running against the same bad input produces the same bad output
Duplicate <code>event_id</code> in source	Not fatal — staging already deduplicates (Assignment 2, Section 3.2); the pipeline succeeds, and <code>staging_duplicate_event_ids</code> stays 0 in the <code>warehouse</code> even though it exists in raw	Handled by design, not by this layer
Disk full / cannot write database file	Transient in principle (space may free up), but capped by the same retry limit as any other transient error so it still escalates to an alert rather than retrying forever	Retrying indefinitely on a persistent condition is itself a reliability bug
Pipeline killed mid-run (e.g. host restart)	No partial-state hazard on the next run	Idempotent full-rebuild design (Assignment 2) means the next scheduled run starts clean regardless of where the previous one stopped
Two runs overlap (previous run still going when the next is scheduled)	Not yet handled	Named explicitly in Section 9 (Limitations) rather than silently ignored — see below
Unclassified / unexpected exception type	Fatal by default, not retried	Avoids masking a real bug behind an automatic retry; surfaces in the alert with a note to add it to the taxonomy

9. Assumptions and Limitations

1. **No overlap/locking guard between scheduled runs.** If a run is still in progress when the next scheduled trigger fires (e.g. an unusually slow run coinciding with the next day’s 2 AM trigger), both could execute against the same SQLite file concurrently. A production version would take a lock file (or a proper job-scheduling system’s built-in overlap prevention) before starting; this coursework version assumes runs complete well within the scheduling interval, consistent with observed run durations (Section 10) of a few seconds against the current data volume.

2. **Retry delay is fixed per run, not adaptive.** `--backoff` supports exponential growth, but nothing currently distinguishes “this particular transient error usually clears in seconds” from “this one usually takes minutes” — every transient error gets the same delay curve.
3. **The Notifier’s real channels (email, Slack) are unconfigured stubs**, per Section 6 — intentional for a credential-free coursework deliverable, not a gap in the interface design.
4. **Monitoring history has no retention policy.** `run_history.jsonl` grows by one line per run indefinitely; a long-running deployment would want to roll it over (mirroring the rotating log handler in Section 3), which this version does not yet do.

10. Observed Behavior

Four scenarios were run against the live pipeline to confirm each path in Section 8 actually behaves as designed (`--simulate` flag, `pipeline_runner.py`):

Simulated condition	Command	Observed outcome
Normal run	<code>pipeline_runner.py</code>	Succeeded, 0 retries, ~3s, all 11 policed checks passed
Transient failure that clears	<code>--simulate transient-then-ok</code>	2 WARNING retries logged, succeeded on 3rd attempt
Persistent transient failure	<code>--simulate always-transient</code> <code>--max-retries 3</code>	3 retries logged, then failed_retries_exhausted , alert fired
Missing source file	<code>--simulate fatal-missing-file</code>	Failed immediately, 0 retries, alert fired, message names the missing file
Corrupted input data	<code>--simulate bad-data</code>	Build succeeded structurally, but <code>run_validation()</code> caught <code>mart_kpi_group_count</code> violating policy; failed immediately, 0 retries, alert names the specific check

All five runs were recorded correctly in `logs/run_history.jsonl`, and `monitor.py` correctly reported the resulting mixed success rate and flagged the most recent failing run — confirming Sections 2 through 6 function as an integrated whole, not just as isolated units.

11. Execution Evidence

The screenshots below are unedited terminal captures from an actual run of this codebase (`assignment3/render_terminal_screenshots.py` renders the raw stdout each command produced into these images — the text is not retyped or paraphrased). They walk through the same scenarios as Section 10, in the order a reviewer would want to verify them: normal operation, both retry outcomes, both non-retried failure types, and the monitoring view that ties all of them together.

11.1 Normal Run

```
assignment3 % python3 pipeline_runner.py

2026-09-11T07:20:44:0000 | INFO | run=20260911T072044Z_f1df8e7c | === Pipeline run 20260911T072044Z_f1df8e7c starting (simulate=None) ===
2026-09-11T07:20:44:0000 | INFO | run=20260911T072044Z_f1df8e7c | [1/5] Building raw layer schema...
2026-09-11T07:20:44:0000 | INFO | run=20260911T072044Z_f1df8e7c | [2/5] Ingesting events CSV into raw_events...
2026-09-11T07:20:46:0000 | INFO | run=20260911T072044Z_f1df8e7c | -> loaded 128974 raw rows
2026-09-11T07:20:46:0000 | INFO | run=20260911T072044Z_f1df8e7c | -> injected 9 deliberately malformed rows (exercises staging validation)
2026-09-11T07:20:46:0000 | INFO | run=20260911T072044Z_f1df8e7c | [3/5] Building staging layer...
2026-09-11T07:20:47:0000 | INFO | run=20260911T072044Z_f1df8e7c | [4/5] Building warehouse layer...
2026-09-11T07:20:47:0000 | INFO | run=20260911T072044Z_f1df8e7c | -> loaded 25 dim_step rows
2026-09-11T07:20:48:0000 | INFO | run=20260911T072044Z_f1df8e7c | [5/5] Building data mart layer...
2026-09-11T07:20:48:0000 | INFO | run=20260911T072044Z_f1df8e7c | Running post-build validation...

--- Data Quality Checks ---
raw_row_count          128983
staging_rejects        6
staging_duplicate_event_ids 0
users_missing_install_event 0
orphan_fact_users      0
completed_without_started 0
activated_without_completed 0
purchased_without_activated 0
invalid_test_group_values 0
mart_kpi_group_count   3
negative_time_to_first_lesson 0
kpi_rates_out_of_bounds 0
2026-09-11T07:20:48:0000 | INFO | run=20260911T072044Z_f1df8e7c | Validation: 11/11 policed checks passed
2026-09-11T07:20:48:0000 | INFO | run=20260911T072044Z_f1df8e7c | === Pipeline run 20260911T072044Z_f1df8e7c SUCCEEDED after 0 retries (4.2s) ===
```

Normal pipeline run: all five layers build, all 11 policed validation checks pass

A complete, unattended run: raw ingestion through data-mart build, the data-quality check output inherited from Assignment 2, and the validation-policy verdict (Section 4) that turns those checks into a pass/fail decision.

11.2 Transient Failure That Recovers

```
assignment3 % python3 pipeline_runner.py --simulate transient-then-ok

2026-09-11T07:20:52:0000 | INFO | run=20260911T072052Z_04342149 | === Pipeline run 20260911T072052Z_04342149 starting (simulate=transient-then-ok) ===
2026-09-11T07:20:52:0000 | WARNING | run=20260911T072052Z_04342149 | Transient failure (attempt 1/5): simulated transient failure (attempt 1) -- retrying in 3s
2026-09-11T07:20:55:0000 | WARNING | run=20260911T072052Z_04342149 | Transient failure (attempt 2/5): simulated transient failure (attempt 2) -- retrying in 3s
2026-09-11T07:20:58:0000 | INFO | run=20260911T072052Z_04342149 | [1/5] Building raw layer schema...
2026-09-11T07:20:58:0000 | INFO | run=20260911T072052Z_04342149 | [2/5] Ingesting events CSV into raw_events...
2026-09-11T07:20:59:0000 | INFO | run=20260911T072052Z_04342149 | -> loaded 128974 raw rows
2026-09-11T07:20:59:0000 | INFO | run=20260911T072052Z_04342149 | -> injected 9 deliberately malformed rows (exercises staging validation)
2026-09-11T07:20:59:0000 | INFO | run=20260911T072052Z_04342149 | [3/5] Building staging layer...
2026-09-11T07:21:00:0000 | INFO | run=20260911T072052Z_04342149 | [4/5] Building warehouse layer...
2026-09-11T07:21:00:0000 | INFO | run=20260911T072052Z_04342149 | -> loaded 25 dim_step rows
2026-09-11T07:21:01:0000 | INFO | run=20260911T072052Z_04342149 | [5/5] Building data mart layer...
2026-09-11T07:21:01:0000 | INFO | run=20260911T072052Z_04342149 | Running post-build validation...

--- Data Quality Checks ---
raw_row_count          128983
staging_rejects        6
staging_duplicate_event_ids 0
users_missing_install_event 0
orphan_fact_users      0
completed_without_started 0
activated_without_completed 0
purchased_without_activated 0
invalid_test_group_values 0
mart_kpi_group_count   3
negative_time_to_first_lesson 0
kpi_rates_out_of_bounds 0
2026-09-11T07:21:01:0000 | INFO | run=20260911T072052Z_04342149 | Validation: 11/11 policed checks passed
2026-09-11T07:21:01:0000 | INFO | run=20260911T072052Z_04342149 | === Pipeline run 20260911T072052Z_04342149 SUCCEEDED after 2 retries (8.9s) ===
```

Two retries logged as WARNING, then the run succeeds on the third attempt

`--simulate transient-then-ok`: the first two attempts raise `TransientError` and are logged as `WARNING`, not `ERROR` — a retry that's still within budget is an expected, non-alarming event. The third attempt succeeds and the run completes normally, with `retries_used=2` recorded in the monitoring history.

11.3 Retries Exhausted → Alert

```
assignment3 % python3 pipeline_runner.py --max-retries 3 --simulate always-transient

2026-09-11T07:21:01+0000 | INFO | run=20260911T072101Z_bf2a0dbd | == Pipeline run 20260911T072101Z_bf2a0dbd starting (simulate=always-transient) ==
2026-09-11T07:21:01+0000 | WARNING | run=20260911T072101Z_bf2a0dbd | Transient failure (attempt 1/3): simulated persistent transient failure (attempt 1) -- retrying in 2s
2026-09-11T07:21:03+0000 | WARNING | run=20260911T072101Z_bf2a0dbd | Transient failure (attempt 2/3): simulated persistent transient failure (attempt 2) -- retrying in 2s
2026-09-11T07:21:05+0000 | WARNING | run=20260911T072101Z_bf2a0dbd | Transient failure (attempt 3/3): simulated persistent transient failure (attempt 3) -- retrying in 2s
2026-09-11T07:21:07+0000 | ERROR | run=20260911T072101Z_bf2a0dbd | Exhausted 3 retries; giving up. Last error: simulated persistent transient failure (attempt 4)
2026-09-11T07:21:07+0000 | CRITICAL | run=- | ALERT [CRITICAL] Pipeline run 20260911T072101Z_bf2a0dbd FAILED after 3 retries :: Last error: simulated persistent transient failure (attempt 4)

Owner/stakeholders should investigate the upstream cause (connection/server issue) before the next scheduled run.

=====
ALERT [CRITICAL]: Pipeline run 20260911T072101Z_bf2a0dbd FAILED after 3 retries
Last error: simulated persistent transient failure (attempt 4)

Owner/stakeholders should investigate the upstream cause (connection/server issue) before the next scheduled run.

=====
2026-09-11T07:21:07+0000 | CRITICAL | run=20260911T072101Z_bf2a0dbd | == Pipeline run 20260911T072101Z_bf2a0dbd FAILED (retries exhausted) ==
```

Three retries, then CRITICAL alert box with the failure reason

`--simulate always-transient --max-retries 3`: three `WARNING`-level retries, then an `ERROR` when the retry budget is exhausted, then the `Notifier`'s `CRITICAL` alert — printed here via the default `LogNotifier` (Section 6), the same call site that would instead page an on-call engineer if `EmailNotifier` or `SlackWebhookNotifier` were configured.

11.4 Fatal Error — No Retry

```
assignment3 % python3 pipeline_runner.py --simulate fatal-missing-file

2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_2b0048b2 | == Pipeline run 20260911T072107Z_2b0048b2 starting (simulate=fatal-missing-file) ==
2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_2b0048b2 | [1/5] Building raw layer schema...
2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_2b0048b2 | [2/5] Ingesting events CSV into raw_events...
2026-09-11T07:21:07+0000 | CRITICAL | run=- | ALERT [CRITICAL] Pipeline run 20260911T072107Z_2b0048b2 FAILED (fatal, not retried) :: source file not found: [Errno 2] No such file or directory: '/home/claude/ab-testing-data-platform/assignment1/generated_data/does_not_exist.csv'. This will not resolve on retry -- check the upstream job that produces does_not_exist.csv, or the --events-csv path, rather than relying on automatic retry.

=====
ALERT [CRITICAL]: Pipeline run 20260911T072107Z_2b0048b2 FAILED (fatal, not retried)
source file not found: [Errno 2] No such file or directory: '/home/claude/ab-testing-data-platform/assignment1/generated_data/does_not_exist.csv'. This will not resolve on retry -- check the upstream job that produces does_not_exist.csv, or the --events-csv path, rather than relying on automatic retry.
=====
2026-09-11T07:21:07+0000 | CRITICAL | run=20260911T072107Z_2b0048b2 | == Pipeline run 20260911T072107Z_2b0048b2 FAILED (fatal) ==
```

Missing source file fails immediately with zero retries

`--simulate fatal-missing-file`: no `WARNING` lines at all — a missing file is classified `FatalError` (Section 2) and goes straight to the alert, since retrying would just repeat an identical failure five times before saying anything useful.

11.5 Validation Failure — No Retry

```
assignment3 % python3 pipeline_runner.py --simulate bad-data

2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_b093c159 | === Pipeline run 20260911T072107Z_b093c159 starting (simulate=bad-data) ===
2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_b093c159 | [1/5] Building raw layer schema...
2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_b093c159 | [2/5] Ingesting events CSV into raw_events...
2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_b093c159 | -> loaded 502 raw rows
2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_b093c159 | -> injected 9 deliberately malformed rows (exercises staging validation)
2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_b093c159 | [3/5] Building staging layer...
2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_b093c159 | [4/5] Building warehouse layer...
2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_b093c159 | -> loaded 25 dim_step rows
2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_b093c159 | [5/5] Building data mart layer...
2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_b093c159 | Running post-build validation...

--- Data Quality Checks ---
raw_row_count          511
staging_rejects        511
staging_duplicate_event_ids 0
users_missing_install_event 0
orphan_fact_users      0
completed_without_started 0
activated_without_completed 0
purchased_without_activated 0
invalid_test_group_values 0
mart_kpi_group_count    0
negative_time_to_first_lesson 0
kpi_rates_out_of_bounds 0

2026-09-11T07:21:07+0000 | INFO | run=20260911T072107Z_b093c159 | Validation: 10/11 policed checks passed
2026-09-11T07:21:07+0000 | CRITICAL | run= | ALERT [CRITICAL] Pipeline run 20260911T072107Z_b093c159 FAILED data validation :: Post-build validation policy violated (n
ot retried -- same input would fail again):
mart_kpi_group_count: value 0 violates policy

=====
ALERT [CRITICAL]: Pipeline run 20260911T072107Z_b093c159 FAILED data validation
Post-build validation policy violated (not retried -- same input would fail again):
mart_kpi_group_count: value 0 violates policy
=====

2026-09-11T07:21:07+0000 | CRITICAL | run=20260911T072107Z_b093c159 | === Pipeline run 20260911T072107Z_b093c159 FAILED (validation) ===
```

Corrupted input data passes ingestion but fails the validation policy

`--simulate bad-data`: this is the scenario that most distinguishes this design from a naive “retry everything” pipeline. The build itself succeeds mechanically — SQL runs, tables populate — but `run_validation()` (Section 4) catches that `mart_kpi_group_count` violates policy and raises `ValidationFailure` before the run can be called successful. The alert names the specific check that failed, not just “something went wrong.”

11.6 Monitoring: Health Summary Across Runs

```
assignment3 % python3 monitor.py --last 15

Pipeline run history -- last 6 of 6 total runs
Run ID      Status      Retries Duration Started
-----
20260911T072044Z_f1df087c success      0      4.22 2026-09-11T07:20:44.285178+00:00
20260911T072052Z_04342149 success      0      0.92 2026-09-11T07:20:52.371414+00:00
20260911T072101Z_bf2a0dbd failed_retries_exhausted 3      6.0 2026-09-11T07:21:01.348525+00:00
20260911T072107Z_2b08048b2 failed_fatal 0      0.01 2026-09-11T07:21:07.385883+00:00
20260911T072107Z_b093c159 failed_validation 0      0.03 2026-09-11T07:21:07.437392+00:00
20260911T072113Z_643fdce1 success      0      3.29 2026-09-11T07:21:13.904126+00:00

Success rate (last 6): 50% Avg duration (successful runs): 5.5s

Most recent run (20260911T072113Z_643fdce1) succeeded.
!/\ Success rate over last 6 runs is below 70% -- investigate even if the most recent run passed.
```

`monitor.py` summarizing six runs, flagging the below-threshold success rate

`monitor.py --last 15` after the five scenarios above plus one final normal run: every outcome — both successes, both failure types, and the retry-exhaustion case — appears as one row each, pulled from `logs/run_history.jsonl`. Because the most recent run happened to succeed, the per-run alert (Section 11.3) would stay quiet — but the success-rate check (Section 5) still flags the run of failures, which is exactly the “pipeline is flapping, not just currently down” case that check exists to catch.

11.7 Scheduler Firing Live

```
assignment3/scheduler % python3 scheduler_demo.py --demo-interval-seconds 4

[scheduler] DEMO MODE: running every 4s (Ctrl+C to stop)

[scheduler] Firing pipeline_runner.py at 2026-09-11 07:21:13
2026-09-11T07:21:13+0000 | INFO | run=20260911T072113Z_643fdce1 | === Pipeline run 20260911T072113Z_643fdce1 starting (simulate=None) ===
2026-09-11T07:21:13+0000 | INFO | run=20260911T072113Z_643fdce1 | [1/5] Building raw layer schema...
2026-09-11T07:21:13+0000 | INFO | run=20260911T072113Z_643fdce1 | [2/5] Ingesting events CSV into raw_events...
2026-09-11T07:21:15+0000 | INFO | run=20260911T072113Z_643fdce1 | -> loaded 128974 raw rows
2026-09-11T07:21:15+0000 | INFO | run=20260911T072113Z_643fdce1 | -> injected 9 deliberately malformed rows (exercises staging validation)
2026-09-11T07:21:15+0000 | INFO | run=20260911T072113Z_643fdce1 | [3/5] Building staging layer...
2026-09-11T07:21:16+0000 | INFO | run=20260911T072113Z_643fdce1 | [4/5] Building warehouse layer...
2026-09-11T07:21:16+0000 | INFO | run=20260911T072113Z_643fdce1 | -> loaded 25 dim step rows
2026-09-11T07:21:17+0000 | INFO | run=20260911T072113Z_643fdce1 | [5/5] Building data mart layer...
2026-09-11T07:21:17+0000 | INFO | run=20260911T072113Z_643fdce1 | Running post-build validation...

--- Data Quality Checks ---
raw_row_count      128983
staging_rejects    6
staging_duplicate_event_ids 0
users_missing_install_event 0
orphan_fact_users  0
completed_without_started 0
activated_without_completed 0
```

`scheduler_demo.py` invoking `pipeline_runner.py` on its own, unattended

`scheduler_demo.py --demo-interval-seconds 4`: the same schedule mechanism as the production crontab entry (Section 7), compressed to a four-second interval so the trigger can be observed directly instead of waiting for a real 2 AM run. The scheduler process itself never touches the pipeline's internals — it only invokes `pipeline_runner.py` as a subprocess, so everything demonstrated in 11.1–11.6 is exactly what fires under the real schedule too.

12. How to Run

```
# Normal run
python pipeline_runner.py

# Demonstrate retry-then-succeed
python pipeline_runner.py --retry-delay-seconds 2 --simulate transient-then-ok

# Demonstrate retries exhausted -> alert
python pipeline_runner.py --retry-delay-seconds 2 --max-retries 3 --simulate always-
transient

# Demonstrate a fatal (non-retried) failure
python pipeline_runner.py --simulate fatal-missing-file
python pipeline_runner.py --simulate bad-data

# Check pipeline health
python monitor.py
python monitor.py --check # exit 1 if unhealthy, for chaining

# Watch the scheduler actually fire, without waiting for 2am
python scheduler/scheduler_demo.py --demo-interval-seconds 15 --retry-delay-seconds 2
# Regenerate the Section 11 evidence screenshots from fresh captures
python render_terminal_screenshots.py
```

See [../diagrams/pipeline_reliability_flow.svg](#) for the full control-flow diagram, and the [repository root README](#) for how this fits into the four-assignment pipeline.