

Assignment 4: Analytics Dashboard & Experiment Evaluation

Onboarding Flow Optimization A/B Test

Aishwarya Mishra | USN 2648610 | MSc Computational Statistics & Applied AI, Christ University

Table of Contents

Abstract

1. KPI Computation

2. Statistical Testing

2.1 Precondition: Sample Ratio Mismatch Check

2.2 Primary Metric: Two-Proportion Z-Tests

2.3 Power Recheck Against Achieved Sample Size

2.4 Secondary Metrics

3. Guardrail Interpretation

4. Recommendation

5. Assumptions and Limitations

6. Dashboard Preview


7. How to Run

Abstract

Assignments 1–3 built the platform; this assignment answers the question that platform exists to answer: did shortening onboarding actually work, and is it worth shipping? This document computes the primary and secondary KPIs from the Assignment 2 data mart, runs the hypothesis test the Assignment 1 experimental design specifies (Section 3), checks the result against every guardrail defined in that same document, and synthesizes a single ship/hold recommendation from all of it — distinguishing, throughout, between what is statistically confirmed and what is merely a directional signal, per the metric-governance model Assignment 1 already committed to.

Analysis Pipeline: Warehouse to Dashboard

Assignment 4 — statistical testing, KPI computation, dashboard generation



Analysis pipeline: warehouse to dashboard

Figure 1. `stats_analysis.py` computes every statistic once, in one place, and writes it to `results.json`; `data_export.py` merges that with the raw mart tables into `dashboard_data.json`; `build_dashboard.py` renders the whole thing into one self-contained HTML file with Plotly.js embedded inline, so `dashboard.html` opens and works fully offline — double-click it, no server, no internet connection required.

1. KPI Computation

Every KPI here is read directly from the Assignment 2 data mart (`mart_kpi_by_group`, `mart_guardrails_by_group`) rather than re-aggregated from raw events — the whole point of that layer (Assignment 2, Section 2.4) was to make this step a lookup, not a re-derivation.

KPI	Value — A	Value — B	Value — C
Onboarding completion rate (primary)	57.4%	73.0%	81.5%
Activation rate	69.8%	63.2%	65.6%
Day-1 return rate	55.8%	54.4%	55.5%

KPI	Value — A	Value — B	Value — C
Day-7 return rate	42.7%	40.1%	42.1%
Day-30 return rate (guardrail)	36.2%	34.3%	29.0%
Paid conversion rate	6.1%	5.7%	6.3%
Avg. time to first lesson	1,584s	1,507s	1,476s

The pattern is exactly what Assignment 1’s opportunity-sizing model (Section 1.3) hoped for on the primary metric — completion rate rises monotonically as step count falls — and exactly what its identification concern (Section 2) warned to watch for on the guardrails: Day-30 return declines as the flow shortens, most sharply for C.

2. Statistical Testing

2.1 Precondition: Sample Ratio Mismatch Check

Before interpreting anything else, `stats_analysis.py` runs the SRM check the Assignment 1 event-tracking plan (Section 3) specifies as mandatory: a chi-square goodness-of-fit test of the realized per-arm allocation against the intended 33/33/33 split.

Result: A=2,000, B=2,000, C=2,000 ($\chi^2=0.0$, $p=1.0$). No SRM detected. The randomizer allocated exactly as intended; every result below can be interpreted at face value.

2.2 Primary Metric: Two-Proportion Z-Tests

The primary hypothesis (Assignment 1, Section 2: H1 — reducing step count increases completion rate) is tested with a two-proportion z-test for each pairwise arm comparison, at $\alpha = 0.05$, with a Wilson-score 95% confidence interval on each arm’s rate and a normal-approximation CI on each pairwise difference.

Comparison	Abs. Diff	Relative Lift	95% CI (diff)	p-value	Result
A vs B	+15.6 pp	+27.2%	[+12.7, +18.5] pp	< 0.0001	Significant
A vs C	+24.1 pp	+42.0%	[+21.3, +26.9] pp	< 0.0001	Significant
B vs C	+8.5 pp	+11.6%	[+5.9, +11.1] pp	< 0.0001	Significant

All three comparisons are significant at $\alpha = 0.05$, and by a wide margin — p-values are not merely below 0.05, they are below floating-point display precision. **H0 is rejected:** onboarding step count has a real, large effect on completion rate within the range tested.

2.3 Power Recheck Against Achieved Sample Size

Assignment 1 (Section 3.4) estimated a design-stage requirement of roughly 1,450–1,500 users per arm to detect a 5-percentage-point difference at 80% power. The achieved sample (2,000 users per arm) is larger than that design target; rechecked against the actual n, this experiment could reliably detect a completion-rate difference as small as **4.3 percentage points** at 80% power. Every observed pairwise difference (8.5–24.1 pp)

is comfortably above that threshold, so none of the three significant results in Section 2.2 is at meaningful risk of being a statistically underpowered false positive.

2.4 Secondary Metrics

Per Assignment 1’s multiple-comparisons reasoning (Section 3.5, Section 4.4), secondary metrics (activation, D1/D7 return, paid conversion, time to first lesson) are reported descriptively and are **not** independently hypothesis-tested — testing five more metrics at the same $\alpha = 0.05$ without correction would inflate the experiment-wise false positive rate, and the primary metric was pre-registered as the sole ship/no-ship basis specifically to avoid that. The one secondary result worth flagging descriptively: **paid conversion rate is not depressed** in B (5.7%) or C (6.3%) relative to A (6.1%) — the central assumption behind Assignment 1’s revenue opportunity model (that checkout mechanics aren’t the bottleneck) holds up in the generated data.

3. Guardrail Interpretation

Per the metric-governance model (Assignment 1, Section 4.4), guardrails are evaluated directionally — flagged as a regression or not, relative to control — rather than hypothesis-tested. Comparing arm C (the aggressive, highest-completion variant) against control:

Guardrail	A (Control)	C (Short)	Flag
Lesson relevance score (1–5)	3.98	3.44	Regression
Refund rate	0.0%	0.0%	No change
Support-ticket rate	0.9%	2.2%	Regression
Notification opt-in rate	52.8%	32.7%	Regression
Day-30 return rate	36.2%	29.0%	Regression (Section 1)

Arm B shows none of these regressions at a comparable magnitude (its worst guardrail movement is a refund-rate increase from 0.0% to 3.8%, on a very small base rate). This is precisely the trade-off Assignment 1’s identification concern (Section 2) predicted might exist: the arm with the *largest* completion-rate win is not the arm with the *cleanest* guardrail profile.

4. Recommendation

SHIP WITH MITIGATION — Arm B (Medium, 7 steps), not Arm C

Arm B beats control by +15.6 percentage points on the primary metric (statistically significant, $p < 0.0001$) with no material guardrail regression. Arm C beats control by more (+24.1 pp) but shows four guardrail regressions relative to control, including the Day-30 return guardrail that exists specifically to catch a shorter flow admitting lower-intent users who complete easily but don’t stick around.

Action: Ship Arm B as the new default onboarding flow. Do not ship Arm C as-is; treat it as a candidate for a follow-up, narrower experiment that investigates whether the specific steps C removes beyond what

B removes (see Assignment 1, Section 3.3's step classification) can be restored selectively, or replaced with a lighter-weight version, to recover more of C's completion-rate lift without C's guardrail cost.

Revenue framing (Assignment 1, Section 1.3 model): applying B's observed +15.6pp completion lift and flat paid-conversion rate to the illustrative 100,000-users/month planning volume implies roughly **+15,600 additional completions/month**, translating to approximately **+890 additional paying users/month** at B's 5.7% conversion rate — a smaller revenue estimate than the full industry-benchmark scenario in Assignment 1, but one that is not purchased with the guardrail cost the Assignment 1 identification concern was specifically designed to catch.

5. Assumptions and Limitations

1. **Simulated data.** As throughout this project (Assignment 1, Section 5), effect sizes come from the calibrated data generator, not observed production users. The statistical *methodology* here (SRM check, two-proportion z-test, power recheck, guardrail governance) is what would be applied identically to real data; the specific numbers in Sections 1–4 are illustrative of that method, not a production forecast.
2. **Guardrails are not corrected for multiple comparisons because they are not hypothesis tests to begin with** (Section 3) — this is a deliberate design choice inherited from Assignment 1, not an oversight; presenting five guardrail p-values would invite exactly the false confidence the governance model is built to avoid.
3. **No heterogeneity/subgroup analysis.** Results are pooled across iOS and Android and across the full install window; a platform-specific or time-windowed (novelty-decay) breakdown is not performed here, though `mart_daily_funnel` (visualized in the dashboard) is available for that follow-up.
4. **The revenue estimate in Section 4 inherits every assumption in Assignment 1's opportunity-sizing model** (illustrative volume and ARPU, flat paid-conversion assumption) — it is a consistent application of that model to this result, not an independent forecast.

6. Dashboard Preview

Onboarding Flow Optimization — Experiment Evaluation Dashboard

3-arm A/B test (13 / 7 / 5 onboarding steps) · Assignment 4 · Aishwarya Mishra, USN 2648610, MSc Computational Statistics & Applied AI, Christ University

Sample Ratio Mismatch Check

Observed allocation: A=2,000, B=2,000, C=2,000 (expected ~2,000 each) — $\chi^2=0.0$, $p=1.0$. No Sample Ratio Mismatch detected -- observed allocation is consistent with the intended 33/33/33 split. Downstream comparisons can be trusted.

Primary Metric: Onboarding Completion Rate

A — Control (13 steps)

57.4%

completion rate · n=2,000
95% CI [55.2%, 59.6%]

B — Medium (7 steps)


73.0%

completion rate · n=2,000
95% CI [71.0%, 74.9%]


C — Short (5 steps)

81.5%

completion rate · n=2,000
95% CI [79.7%, 83.1%]



Dashboard top: SRM check, KPI cards, and the completion-rate chart with 95% confidence intervals



Recommendation

SHIP WITH MITIGATION (Arm B, not C)

B significantly beats control on completion rate with no major guardrail regression. C wins by more on completion but shows 3 guardrail regressions relative to control -- ship B now; treat C as a follow-up experiment after addressing what's driving its guardrail declines.

Generated from assignment2/db/ab_test.db via assignment4/analysis/stats_analysis.py and assignment4/dashboard/data_export.py · Full methodology in 00_documentation.md

Dashboard bottom: funnel diagnostics and the final ship/hold recommendation

7. How to Run

```
# 1. Compute statistics from the Assignment 2 warehouse
cd assignment4/analysis
python stats_analysis.py --db ../../assignment2/db/ab_test.db

# 2. Export dashboard data (mart tables + stats results)
cd ../dashboard
python data_export.py --db ../../assignment2/db/ab_test.db

# 3. Build the dashboard
python build_dashboard.py

# Open dashboard/dashboard.html directly in any browser -- no server needed
```

See [dashboard/dashboard.html](#) for the interactive dashboard and [executive_summary.pdf](#) for a one-page presentation-style summary of this document.